

1. Importing necessary Modules

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
```

2. Uploading the given Dataset

```
In [2]: data=pd.read_csv("abalone.csv")
data.head()
```

```
Out[2]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

```
In [3]: data.shape
```

```
Out[3]: (4177, 9)
```

```
In [4]: #Modifying the given dataset
Age=1.5*data.Rings
data["Age"]=Age
data=data.rename(columns = {'Whole weight': 'Whole_weight', 'Shucked weight': 'Shucked_weight', 'Viscera weight': 'Viscera_weight',
'Shell weight': 'Shell_weight'})
data=data.drop(columns=["Rings"],axis=1)
data.head()
```

```
Out[4]:
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

```
In [5]: data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sex              4177 non-null   object
1   Length           4177 non-null   float64
2   Diameter         4177 non-null   float64
3   Height           4177 non-null   float64
4   Whole_weight     4177 non-null   float64
5   Shucked_weight   4177 non-null   float64
6   Viscera_weight   4177 non-null   float64
7   Shell_weight     4177 non-null   float64
8   Age              4177 non-null   float64
dtypes: float64(8), object(1)
memory usage: 293.8+ KB
```

3.Performing various Visualizations

Univariate Analysis


```
In [6]: sns.histplot(data["Age"],color='blue')
```

```
Out[6]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5db9fa2690>
```



```
In [7]: sns.displot(data["Age"],color='yellow')
```


```
Out[7]: <seaborn.axisgrid.FacetGrid at 0x7f5db9e01a10>
```



```
In [8]: sns.boxplot(data["Age"],color='red')
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
FutureWarning
```

```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5db70117d0>
```




Bivariate Analysis

```
In [9]: sns.barplot(data["Height"],data["Whole_weight"])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
FutureWarning
```


```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5db6f99750>
```



```
In [10]: sns.lineplot(data["Height"],data["Whole_weight"], color='blue')
```


```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
FutureWarning
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5db6cfcdb0>
```



```
In [11]: sns.scatterplot(x=data.Height,y=data.Whole_weight,color='green')
```


```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5db6d38110>
```



```
In [12]: sns.regplot(data['Height'],data['Whole_weight'],color='pink')
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

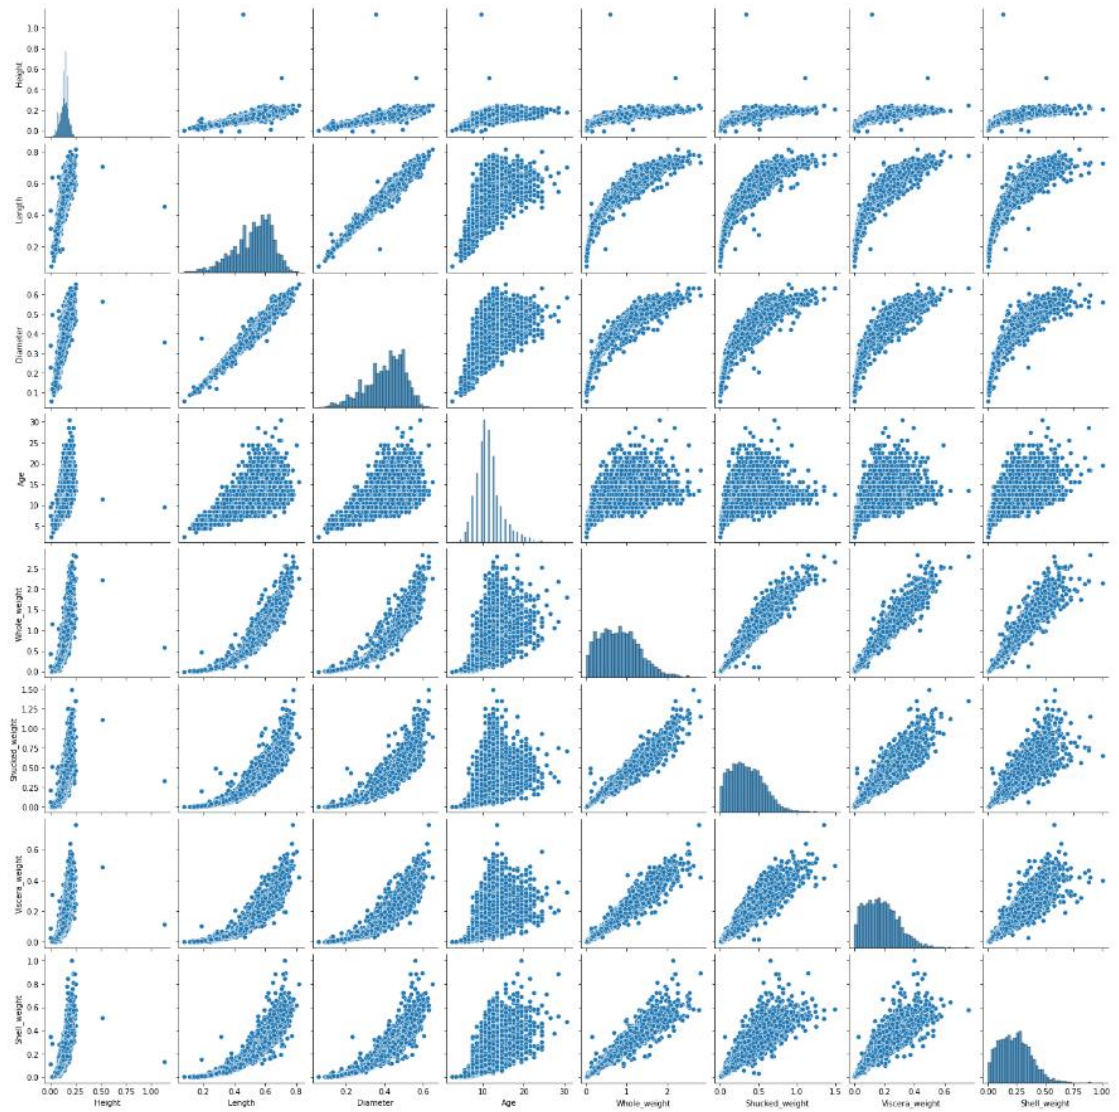
```
FutureWarning
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5db6d5f9d0>
```



Multivariate Analysis

```
In [13]: sns.pairplot(data=data[["Height","Length","Diameter","Age","Whole_weight","Shucked_weight","Viscera_weight","Shell_weight"]])
```

```
Out[13]: <seaborn.axisgrid.PairGrid at 0x7f5db6c59d10>
```



4. Performing descriptive statistics on the dataset.

```
In [14]: data.describe(include='all')
```

```
Out[14]:
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
count	4177	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
unique	3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
top	M	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
freq	1528	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
mean	NaN	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	11.433684
std	NaN	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
min	NaN	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	2.500000
25%	NaN	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	9.500000
50%	NaN	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	10.500000
75%	NaN	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	12.500000
max	NaN	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	30.500000

5. Checking for Missing values

```
In [15]: data.isnull().sum()
```

```
Out[15]: Sex          0
Length          0
Diameter        0
Height          0
Whole_weight    0
Shucked_weight  0
Viscera_weight  0
Shell_weight    0
Age            0
dtype: int64
```

6. Finding the outliers and replacing them outliers

```
In [16]: outliers=data.quantile(q=(0.25,0.75))
outliers
```

```
Out[16]:
```


	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0.25	0.450	0.35	0.115	0.4415	0.186	0.0935	0.130	9.5
0.75	0.615	0.48	0.165	1.1530	0.502	0.2530	0.329	12.5

```
In [17]: q1 = data.Age.quantile(0.25)
q3 = data.Age.quantile(0.75)
IQR = q3 - q1
lower_limit = q1 - 1.5 * IQR
data.median(numeric_only=True)
```

```
Out[17]: Length      0.5450
Diameter    0.4250
Height      0.1400
Whole_weight 0.7995
Shucked_weight 0.3360
Viscera_weight 0.1710
Shell_weight 0.2340
Age         10.5000
dtype: float64
```

```
In [18]: data['Age'] = np.where(data['Age'] < lower_limit, 7, data['Age'])
sns.boxplot(x=data.Age, showfliers = False)
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5db4d71950>
```



7.Checking for Categorical columns and perform Encoding

```
In [19]: data.head()
```

```
Out[19]:
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

8. Split the data into dependent and independent variables

```
In [22]: y = data["Sex"]  
y.head()
```

```
Out[22]: 0    2  
1    2  
2    0  
3    2  
4    1  
Name: Sex, dtype: int64
```

```
In [23]: x=data.drop(columns=["Sex"],axis=1)  
x.head()
```

```
Out[23]:
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

9. Scaling the values

```
In [24]: from sklearn.preprocessing import scale  
X_Scaled = pd.DataFrame(scale(x), columns=x.columns)  
X_Scaled.head()
```

```
Out[24]:
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	-0.574558	-0.432149	-1.064424	-0.641898	-0.607685	-0.726212	-0.638217	1.577830
1	-1.448986	-1.439929	-1.183978	-1.230277	-1.170910	-1.205221	-1.212987	-0.919022
2	0.050033	0.122130	-0.107991	-0.309469	-0.463500	-0.356690	-0.207139	-0.294809
3	-0.699476	-0.432149	-0.347099	-0.637819	-0.648238	-0.607600	-0.602294	0.017298
4	-1.615544	-1.540707	-1.423087	-1.272086	-1.215968	-1.287337	-1.320757	-0.919022

10. Split the data into training and testing

```
In [25]: from sklearn.model_selection import train_test_split
X_Train, X_Test, Y_Train, Y_Test = train_test_split(X_Scaled, y, test_size=0.2, random_state=0)
```

```
In [26]: X_Train.shape, X_Test.shape
```

```
Out[26]: ((3341, 8), (836, 8))
```

```
In [27]: Y_Train.shape, Y_Test.shape
```

```
Out[27]: ((3341,), (836,))
```

```
In [28]: X_Train.head()
```

```
Out[28]:
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
3141	-2.864726	-2.750043	-1.423087	-1.622870	-1.553902	-1.583867	-1.644065	-1.543234
3521	-2.573250	-2.598876	-2.020857	-1.606554	-1.551650	-1.565619	-1.626104	-1.387181
883	1.132658	1.230689	0.728888	1.145672	1.041436	0.286552	1.538726	1.577830
3627	1.590691	1.180300	1.446213	2.164373	2.661269	2.330326	1.377072	0.017298
2106	0.591345	0.474853	0.370226	0.432887	0.255175	0.272866	0.906479	1.265723

```
In [29]: X_Test.head()
```

```
Out[29]:
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
668	0.216591	0.172519	0.370226	0.181016	-0.368878	0.569396	0.690940	0.953617
1580	-0.199803	-0.079426	-0.466653	-0.433875	-0.443224	-0.343004	-0.325685	-0.606915
3784	0.799543	0.726798	0.370226	0.870348	0.755318	1.764639	0.565209	0.329404
463	-2.531611	-2.447709	-2.020857	-1.579022	-1.522362	-1.538247	-1.572219	-1.543234
2615	1.007740	0.928354	0.848442	1.390405	1.415417	1.778325	0.996287	0.641511

```
In [30]: Y_Train.head()
```

```
Out[30]:
3141    1
3521    1
883     2
3627    2
2106    2
Name: Sex, dtype: int64
```

```
In [31]: Y_Test.head()
```

```
Out[31]:
668     2
1580    1
3784    2
463     1
2615    2
Name: Sex, dtype: int64
```

11. Build the Model

```
In [32]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(n_estimators=10, criterion='entropy')
```

```
In [33]: model.fit(X_Train, Y_Train)
```

```
Out[33]: RandomForestClassifier(criterion='entropy', n_estimators=10)
```

```
In [34]: y_predict = model.predict(X_Test)
```

```
In [35]: y_predict_train = model.predict(X_Train)
```

12. Train the Model

```
In [36]: from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [37]: print('Training accuracy: ', accuracy_score(Y_Train, y_predict_train))
```

```
Training accuracy: 0.9811433702484286
```

13. Test the Model

```
In [38]: print('Testing accuracy: ', accuracy_score(Y_Test, y_predict))
```

Testing accuracy: 0.5526315789473685

14. Measure the performance using Metrics

```
In [39]: pd.crosstab(Y_Test, y_predict)
```

```
Out[39]: col_0  0   1   2
Sex
0  112  33  104
1   45 211   35
2   100  57  139
```

```
In [40]: print(classification_report(Y_Test, y_predict))
```

	precision	recall	f1-score	support
0	0.44	0.45	0.44	249
1	0.70	0.73	0.71	291
2	0.50	0.47	0.48	296
accuracy			0.55	836
macro avg	0.55	0.55	0.55	836
weighted avg	0.55	0.55	0.55	836